


```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
%matplotlib inline

df = pd.read_csv('/content/height-weight.csv')
```


```
#scatter plot
plt.scatter(df['Weight'], df['Height'])
plt.xlabel("Weight")
plt.ylabel("Height")
```

```
Text(0, 0.5, 'Height')
```




```
#seaborn for visualization
import seaborn as sns
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7f08da5d5cc0>
```




```
x = df[['Weight']]
y = df[['Height']]
```

```
X_series=df['Weight']  
np.array(X_series).shape
```

 (23,)

```
np.array(y).shape
```

 (23, 1)

```
## Train Test Split  
from sklearn.model_selection import train_test_split
```


```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.25, random_state=42)
```

```
#standardization  
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```


```
X_test
```

 array([[0.33497168],
[0.33497168],
[-1.6641678],
[1.36483141],
[-0.45256812],
[1.97063125]])

```
from sklearn.linear_model import LinearRegression
```


```
regression=LinearRegression(n_jobs=-1)
```

```
regression.fit(X_train, y_train)
```





LinearRegression
LinearRegression(n_jobs=-1)

```
print("Coefficient or slope:", regression.coef_)  
print("Intercept:", regression.intercept_)
```

 Coefficient or slope: [[17.2982057]]
Intercept: [156.47058824]

```
plt.scatter(X_train, y_train)  
plt.plot(X_train, regression.predict(X_train))
```

 [`<matplotlib.lines.Line2D at 0x7f08cf764bb0>`]




Start coding or [generate](#) with AI.

prediction of test data predicted height output = intercept + coef(Weights) $y_{pred_test} = 156.470 + 17.29(X_{test})$

```
#prediction of test data
y_pred = regression.predict(X_test)
```

```
#performance Metrics
from sklearn.metrics import mean_absolute_error, mean_squared_error
```


```
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print(mae)
print(mse)
print(rmse)
```

 9.66512588679501
114.84069295228699
10.716374991212605

R square formula $R^2 = 1 - \frac{SSR}{SST}$ R^2 = coefficient of determination SSR = sum of squares of residuals SST = total sum of squares

```
from sklearn.metrics import r2_score
```


```
score = r2_score(y_test, y_pred)
print(score)
```

 0.7360826717981276


```
##Ols Linear Regression
import statsmodels.api as sm
```

```
model = sm.OLS(y_train, X_train).fit()
```

```
prediction = model.predict(X_test)
print(prediction)
```

 [5.79440897 5.79440897 -28.78711691 23.60913442 -7.82861638
34.08838469]

```
print(model.summary())
```

 OLS Regression Results
=====

```
Dep. Variable:      Height    R-squared (uncentered):      0.012
Model:              OLS      Adj. R-squared (uncentered): -0.050
Method:             Least Squares  F-statistic:      0.1953
Date:               Mon, 22 Jul 2024  Prob (F-statistic):    0.664
Time:               07:48:56   Log-Likelihood:   -110.03
No. Observations:   17        AIC:      222.1
Df Residuals:       16        BIC:      222.9
Df Model:            1
Covariance Type:    nonrobust
```

	coef	std err	t	P> t	[0.025	0.975]
x1	17.2982	39.138	0.442	0.664	-65.671	100.267
Omnibus:	0.135		Durbin-Watson:		0.002	
Prob(Omnibus):	0.935		Jarque-Bera (JB):		0.203	
Skew:	-0.166		Prob(JB):		0.904	
Kurtosis:	2.581		Cond. No.		1.00	

Notes:

[1] R² is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

/usr/local/lib/python3.10/dist-packages/scipy/stats/_stats_py.py:1806: UserWarning: kurtosistest only valid for n>=20 ... continuing an warnings.warn("kurtosistest only valid for n>=20 ... continuing ")

```
#prediction for new data
regression.predict(scaler.transform([[72]]))

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was f
warnings.warn(
array([[155.97744705]])
```

Start coding or [generate](#) with AI.